


FEED FORWARD NEURAL NETWORK

- Amount of nodes ①

Making a feedforward neural network which will classify handwritten digit images from MNIST dataset.

Architecture of our neural net will be something like this,



This is a 4 layered neural net, input will be (784,1), layer 1 will be (128,1), layer 2 will be (128,1) and output layer will be (10,1) for all digits from 0 to 9.

Let's look at the dataset. We'll use Kaggle's MNIST dataset.

labels	pixel 0	pixel 1	pixel 2	...	pixel 783
5	0	0	255		0
9
3
4
1

defining our 'network' ~~size~~ class.

```
class Network(object):
```

```
    def __init__(self, sizes):
```

```
        self.num_layers = num(sizes)
```

```
        self.sizes = sizes
```

```
        self.biases = [np.random.randn(y) for y in sizes[1:]]
```

```
        self.weights = [np.random.randn(y, x) for x, y in
```

```
            zip(sizes[::-1], sizes[1:])]

object = (784, 30, 10)
```

```
biases = [  $\begin{bmatrix} b_1 \\ (30, 1) \end{bmatrix}$  ,  $\begin{bmatrix} b_2 \\ (10, 1) \end{bmatrix}$  ] weights = [  $\begin{bmatrix} w_1 \\ (30, 784) \end{bmatrix}$  ,  $\begin{bmatrix} w_2 \\ (10, 30) \end{bmatrix}$  ]
```

```
def feedforward(self, a):
```

```
    for bias in zip(self.biases, self.weights):
```

```
        a = sigmoid(np.dot(wia) + b)
```

```
    return a
```

In this function, a will be our input feature of dimension, $(784, 1)$. for our current values this loop will go two times for (b_1, w_1) and (b_2, w_2)

$$a_1 = \text{sigmoid}(\text{np.dot}(w_1, a_0) + b_1) \rightarrow (30, 1)$$

$$a_2 = \text{sigmoid}(\text{np.dot}(w_2, a_1) + b_2) \rightarrow (10, 1)$$

$$w_1 \equiv (30, 784)$$

$$a_0 \equiv (784, 1)$$

$$b_1 \equiv (30, 1)$$

$$w_2 \equiv (10, 30)$$

$$a_1 \equiv (30, 1)$$

$$b_2 \equiv (10, 1)$$

Gradient descent

```
for _ in range(epochs):
```

```
    nabl_a_b = [np.zeros(b.shape) for b in self.biases]
```

```
    nabl_a_w = [np.zeros(w.shape) for w in self.weights]
```

```
    for i in range(m):
```

```
        xx, yy = x[i], y[i]
```

```
        delta_nabl_a_b, delta_nabl_a_w = self.backprop(xx, yy)
```

```
        nabl_a_b = [nb + dnb for nb, dnb in zip(nabl_a_b, delta_nabl_a_b)]
```

```
        nabl_a_w = [nw + dnw for nw, dnw in zip(nabl_a_w, delta_nabl_a_w)]
```

```
    self.weights = [w - (alpha/m) * nw for w, nw in zip(self.weights, nabl_a_w)]
```

```
    self.biases = [b - (alpha/m) * nb for b, nb in zip(self.biases, nabl_a_b)]
```

So in gradient descent, we'll start with matrices of 0 like $\text{nabla}_a \mathbf{b}$ and $\text{nabla}_a \mathbf{w}$ for each epoch.

Inside each epoch,

for each example one by one,

we'll calculate δ i.e. $\text{delta_nabla}_a \mathbf{b}$ and $\text{delta_nabla}_a \mathbf{w}$ for each example and ~~append~~ ^{add} it to our previously 0 initialized vectors $\text{nabla}_a \mathbf{b}$ and $\text{nabla}_a \mathbf{w}$.

Now we have matrices $\text{nabla}_a \mathbf{w}$ and $\text{nabla}_a \mathbf{b}$ which have weights and biases ~~changed~~ ^{added} delta values for all of the examples of that epoch.

Now we'll use this $\nabla \mathbf{w}$ and $\nabla \mathbf{b}$ to make changes to our weights.

Backpropagation.

Backpropagation algorithm deals with one example at a time. And the cost function which we'll use here is a simple mean-squared error function instead of the more commonly used cross-entropy error function.

$$C = \frac{1}{2} \|\mathbf{y} - \mathbf{a}^L\|^2 = \frac{1}{2} \sum_j (\mathbf{y}_j - \mathbf{a}_j^L)^2$$

Equations of backpropagation

Backprop is about understanding how changing the weights and biases in a network changes the cost function. Ultimately this means computing the partial derivative $\frac{\partial C}{\partial \mathbf{w}_{jk}^L}$ and $\frac{\partial C}{\partial \mathbf{b}_j^L}$. But to compute those we'll first introduce a intermediate quantity (δ_j^L) delta, which we can call as the error in the j^{th} neuron of layer L .

Backpropagation will give us a procedure to compute δ_j^L and then will relate (δ_j^L) to $\left(\frac{\partial C}{\partial \mathbf{w}_{jk}^L}\right)$ and $\left(\frac{\partial C}{\partial \mathbf{b}_j^L}\right)$.

Defining the error in a neuron

$$\delta_j^L = \frac{\partial C}{\partial z_j^L}$$



$$\delta_j^L = \frac{\partial C}{\partial z_j^L} \equiv \text{how fast the cost is changing with } (z_j^L).$$

$$\delta_j^L = \frac{\partial C}{\partial z_j^L} = \left(\frac{\partial C}{\partial a_j^L} \right) \cdot \left(\frac{\partial a_j^L}{\partial z_j^L} \right) \equiv \begin{matrix} \text{(how fast cost is changing} \\ \text{with activation)} \\ \text{(how fast* activation} \\ \text{is changing with } z_j^L) \end{matrix}$$

but

$$a_j^L = \sigma(z_j^L), \text{ hence}$$

$$\frac{\partial}{\partial z_j^L} (\sigma(z_j^L)) = \sigma'(z_j^L)$$

∴

$$\boxed{\delta_j^L = \frac{\partial C}{\partial z_j^L} = \left(\frac{\partial C}{\partial a_j^L} \right) \cdot \sigma'(z_j^L)} \quad \text{--- ①}$$

equation for the error in output layer.

The derivative of the sigmoid function will be $g(z) * (1 - g(z))$ where $g(z)$ is the sigmoid function.

Derivative of the cost function we have chosen will be simply $(a^L - y)$.

∴ the vectorized implementation of equation ① will be

$$\delta^L = \nabla_a C \odot \sigma'(z^L)$$

$$\boxed{\delta^L = (a^L - y) \odot \sigma'(z^L)} \quad \text{--- ② vectorised.}$$

An equation for the error δ^L in terms of the error in the next layer, δ^{L+1} ,

$$\delta_j^L = \frac{\partial C}{\partial z_j^L}$$

similarly

$$\delta_j^{L+1} = \frac{\partial C}{\partial z_j^{L+1}}$$



$\left(\delta^L = \frac{\partial C}{\partial z^L} \right)$, now by using chain rule this could be written as,

$$\delta^L = \left(\frac{\partial C}{\partial z^{L+1}} \right) \left(\frac{\partial z^{L+1}}{\partial z^L} \right)$$

$$\delta_L = \left(\frac{\partial z^{L+1}}{\partial z^L} \right) \delta^{L+1} \quad \text{--- (i)}$$

$$z^{L+1} = (w^{L+1}) \cdot a^L + b^{L+1} \quad \text{--- (ii)}$$

for combining (i) and (ii):

$$\frac{\partial z^{L+1}}{\partial z^L} = \frac{\partial}{\partial z^L} (w^{L+1} \cdot \sigma(z^L) + b^{L+1})$$

$$\frac{\partial z^{L+1}}{\partial z^L} = w^{L+1} \cdot \sigma'(z^L)$$

and hence,

$$\delta_L = w^{L+1} \cdot \sigma'(z^L) \times \delta^{L+1} \quad \text{--- (2)}$$

$$\delta^L = ((w^{L+1})^T \cdot \delta^{L+1}) \odot \sigma'(z^L) \quad \text{--- (2) vectorized.}$$

How does cost of the network changes with the bias? (6)

It turns out,

$$\boxed{\frac{\partial C}{\partial b_j^L} = \delta_j^L} \text{ --- (3)}$$

$$\boxed{\frac{\partial C}{\partial b} = \delta} \rightarrow \text{vectorized (3)}$$

Now, similarly we want to find a way to compute change in cost w.r.t. change in weight of a particular neuron.

Equation (1) gives us a way of finding change in cost with δ .

Equation (2) finds a way to relate weights ~~with~~ of next layer with δ .

$$\boxed{\delta^L = \nabla_a C \odot \sigma'(z^L)} \text{ --- (1)}$$

$$\boxed{\delta^L = ((w^{L+1})^T \delta^{L+1}) \odot \sigma'(z^L)} \text{ --- (2)}$$

Now we'll do some basic algebra to reach a rough conclusion to the required relation,

$$\nabla_a C \odot \sigma'(z^L) = (w^{(L+1)T} \cdot \delta^{L+1}) \odot \sigma'(z^L)$$


$$\cancel{w^{L+1}} \cdot \cancel{\delta^{L+1}} = \cancel{\nabla_a C} \cdot \cancel{\sigma'(z^L)}$$

$$\nabla_a C = \boxed{\frac{\partial C}{\partial a^L} = (w^{L+1}) \cdot \delta^{L+1}}$$

$$\boxed{\frac{\partial C}{\partial w^{L+1}} = a^L \cdot \delta^{L+1}} \rightarrow \text{now replacing } L \text{ by } L-1$$

$$\boxed{\frac{\partial C}{\partial w^L} = a^{L-1} \cdot \delta^L} \text{ --- (4)}$$

Intuitive meaning of equation (9) -



$$\boxed{\frac{\partial C}{\partial w} = a^{L-1} \cdot \delta^L}$$

Code for Backpropagation.

```
def backprop(self, x, y):  
    nabla_b = [np.zeros(b.shape) for b in self.biases]  
    nabla_w = [np.zeros(w.shape) for w in self.weights]  
  
    ## feedforward propagation.  
    x = x.reshape(-1, 1)  ## adjusting rank 1 python arrays.  
    activation = x  
    activations = [x]  
    zs = []  
  
    for b, w in zip(self.biases, self.weights):  
        z = np.dot(w, activation) + b  
        zs.append(z)  
        activation = sigmoid(z)  
        activations.append(activation)
```

/* Above code does all the feedforward work and stores all the z's and the activations for later use during backward pass step */

backward propagation.

delta = self.cost_derivative(activations[-1], y) * sigmoid_derivative(zs[-1])

↓ calculating δ for the first layer

nabla_b[-1] = delta

nabla_w[-1] = np.dot(delta, activation[-2].T)

for l in range(2, self.num_layers):

z = zs[l]

sp = sigmoid_derivative(z)

delta = np.dot(self.weights[l+1].T, δ) * sp

nabla_b[l] = delta

nabla_w[l] = np.dot(delta, activations[l+1].T)

return(nabla_b, nabla_w)

/* In the above code, all we did was to implement the equations of backpropagation for all the layers and we made an efficient use of negative indices of python */

Cost Derivative.

for our mean squared cost function, $\frac{1}{2}(a-y)^2$ its derivative suitable here will be, $2(a-y)$ or we can simply use $(a-y)$.

code.

derivative of cost function.

def cost_derivative(self, output_activations, y):

k = np.zeros((10,1))

k[y] = 1

return (output_activation - k)